

Literature survey on
Decoding with language models

CS-E4070 End-to-end systems for speech recognition

Anssi Moisio, 474694

May 16, 2020

1 Introduction

An End-to-End (E2E) automatic speech recognition (ASR) system is trained on a parallel corpus of speech and transcription of the speech. Current E2E ASR systems, which are typically neural network models, require large numbers of data to train on. This poses a problem because transcribed speech data are arduous to produce and therefore expensive. The problem becomes crucial especially when training E2E ASR systems for low-resource languages that have few data to train the models on. In contrast to parallel corpora, text-only corpora are usually abundantly available even for low-resource languages. Therefore a proposed solution to the problem is to use text data to learn a language model, which can then be integrated into the E2E ASR system to provide information about the language at the decoding stage. This also enables domain transfer learning: an E2E ASR system can be trained on a source domain in which there are large parallel corpora available, and then transferred to a target domain by integrating a target domain language model to the system.

2 Background

2.1 Decoding with language models in traditional ASR

Traditional ASR systems are influenced by the noisy channel model framework from information theory (Jelinek, 1976). In this framework, a *source* (speaker) produces a piece of text. The task of a language model (LM) is to estimate the prior probability $P_{LM}(\mathbf{y})$ that a piece of text \mathbf{y} is produced by the source. The piece of text then goes through a *channel* that encodes it and transforms it into a sequence of phonemes \mathbf{x} which can be mapped to graphemes using a phoneme lexicon (also called dictionary). The channel is modelled by an acoustic model that generates a probability distribution

$$P(\mathbf{x}|\mathbf{y}) = \sum_{s \in S_{\mathbf{y}}} P(\mathbf{x}|s)P(s|\mathbf{y}) = \sum_{s \in S_{\mathbf{y}}} \prod_t P(x_t|s(t)) \quad (1)$$

for acoustic features $\mathbf{x} = \mathbf{x}_1, \dots, \mathbf{x}_T$ given a piece of text $\mathbf{y} = y_1, \dots, y_U$, with possible time alignments $S_{\mathbf{y}} = \{\dots, s, \dots\}$ (McDermott et al., 2019). The prior alignments $P(s|\mathbf{y})$ can be implemented using, for example, a Markov model. A *decoder* then determines the piece of text that has most probably caused the observed sequence, i.e. the output of the acoustic model mapped to graphemes. Using the Bayes' formula, this means maximising the posterior probability

$$P(\mathbf{y}|\mathbf{x}) = \frac{P(\mathbf{x}|\mathbf{y})P_{LM}(\mathbf{y})}{P(\mathbf{x})} \quad (2)$$

where the probability of the acoustic features $P(\mathbf{x})$ can be ignored because it does not affect the maximum posterior, yielding the maximum posterior estimate

$$\hat{\mathbf{y}} = \underset{\mathbf{y}}{\operatorname{argmax}} P(\mathbf{x}|\mathbf{y})P_{LM}(\mathbf{y}) \quad (3)$$

This approach works well for traditional ASR systems that use generative models for acoustic modelling, for example gaussian mixture models. However, if the generative model is replaced with a discriminative neural network model, it does not produce an acoustic likelihood $P(\mathbf{x}_t|\mathbf{s}(t))$ but a state level posterior probability $P(\mathbf{s}(t)|\mathbf{x}_t)$, which is a problem because the decoding (eq. 3) relies on the likelihoods. This mismatch can be bypassed by applying the Bayes' rule to produce pseudo-likelihoods:

$$P(\mathbf{x}_t|\mathbf{s}(t)) = \frac{P(\mathbf{s}(t)|\mathbf{x}_t)P(\mathbf{x}_t)}{P(\mathbf{s}(t))} \propto \frac{P(\mathbf{s}(t)|\mathbf{x}_t)}{P(\mathbf{s}(t))} \quad (4)$$

The state priors $P(\mathbf{s}(t))$ can be gathered from corpus frequencies (Bourlard and Morgan, 2012).

2.2 Seq2Seq encoder-decoder for ASR

E2E ASR systems map the input sequence \mathbf{x} directly to text sequence \mathbf{y} without decoupling the language model, the acoustic model and the phoneme to grapheme mapping. A common E2E speech recognition system architecture is that of a sequence-to-sequence (Seq2Seq) (Sutskever et al., 2014) encoder-decoder. A Seq2Seq model maps an input sequence to an output sequence. In ASR, the input sequence is typically acoustic features extracted from a speech signal and the output sequence is typically characters, sub-word units, or words. The encoder creates a hidden representation \mathbf{h} of the input sequence, and the decoder uses this representation to generate the probabilities of the output sequences

$$P(\mathbf{y}|\mathbf{x}) = P(\mathbf{y}|\mathbf{h}) = \prod_t P(y_t|\mathbf{h}, \mathbf{y}_{<t}) \quad (5)$$

Attention mechanisms, as introduced by Bahdanau et al. (2014), can be used between the encoder and the decoder to learn to locate and use the most relevant parts of the encoded hidden representation for the decoder to generate a certain part of the output sequence. An influential encoder-decoder model with attention for ASR was published by Chan et al. (2016), called Listen, Attend and spell (LAS). In this model, the attention mechanism calculates a context vector $\mathbf{c}_t = \text{AttentionContext}(\mathbf{d}_t, \mathbf{h})$ for each time step. The current decoder state \mathbf{d}_t is determined by the previous context vector \mathbf{c}_{t-1} , previous decoder state \mathbf{d}_{t-1} , and the previously emitted output $\hat{\mathbf{y}}_{t-1}$

$$\mathbf{d}_t = \text{RNN}(\hat{\mathbf{y}}_{t-1}, \mathbf{d}_{t-1}, \mathbf{c}_{t-1}) \quad (6)$$

where RNN is a 2 layer LSTM (Long Short-Term Memory) network. The decoder outputs the posterior distribution at a time step t after a final layer

$$P(y_t|\mathbf{h}, \mathbf{y}_{<t}) = \text{softmax}(\mathbf{W}_s[\mathbf{d}_t, \mathbf{c}_t] + \mathbf{b}_s) \quad (7)$$

where \mathbf{W}_s and \mathbf{b}_s are learnable parameters.

3 LM integration approaches

3.1 Shallow fusion and density ratio approach

In the LM integration approach that (Gulcehre et al., 2015) termed *shallow fusion* in contrast to *deep fusion*, the LM and ASR model output scores are combined to compute the output sequence similarly to the traditional ASR systems (eq. 3):

$$\hat{\mathbf{y}} = \underset{\mathbf{y}}{\operatorname{argmax}} \log P(\mathbf{y}|\mathbf{x}) + \lambda \log P_{LM}(\mathbf{y}) \quad (8)$$

where λ is a tunable language model weight, and the logarithmic probabilities are used to avoid numeric underflow.

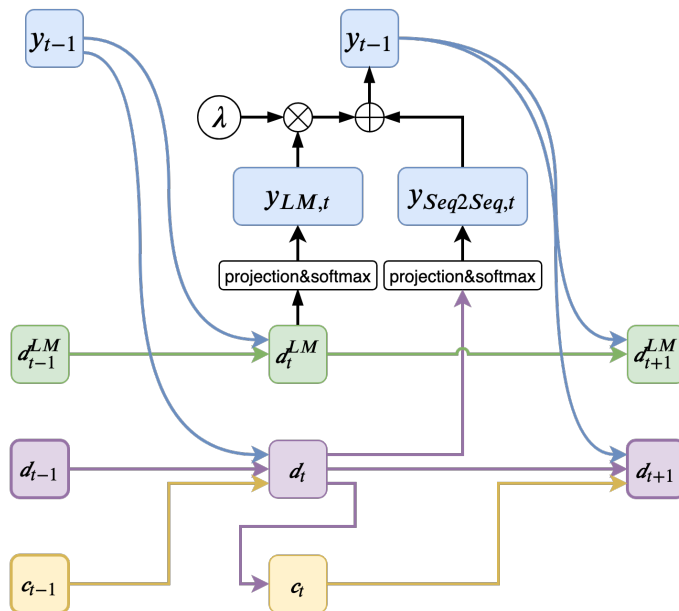


Figure 1: Shallow fusion decoding step. The LM and Seq2Seq model outputs are combined.

In practice, $\log P(\mathbf{y}|\mathbf{x})$ is typically approximately maximized using a left-to-right beam search algorithm that produces a list of K candidates for \mathbf{y} . The combined score $\log P(\mathbf{y}|\mathbf{x}) +$

$\lambda \log P_{LM}(\mathbf{y})$ is calculated only for K candidate hypotheses $\{\hat{\mathbf{y}}^{(i)}\}_{i=1,\dots,K}$ with the top scores given by the decoder from the complete hypothesis space.

The shallow fusion approach can be extended to a density ratio approach if we have language models for both the source and target domain (McDermott et al., 2019). In this approach, the Seq2Seq probability is combined with the density ratio of the target and source domain probability distributions, instead of the single LM probability distribution as in shallow fusion. The division by the source domain LM output is a subtraction in the log domain:

$$\hat{\mathbf{y}} = \underset{\mathbf{y}}{\operatorname{argmax}} \log P(\mathbf{y}|\mathbf{x}) + \lambda_{target} \log P_{LM,target}(\mathbf{y}) - \lambda_{source} \log P_{LM,source}(\mathbf{y}) \quad (9)$$

Here the Seq2Seq model is trained on the source domain parallel (audio, transcript) corpus and the source domain LM is trained on the same transcripts.

3.2 Deep fusion

In shallow fusion, the LM is weighted uniformly with the scalar λ across all hypotheses. Intuitively, however, the language model is more useful when decoding some sequences than others, and should therefore be weighted depending on the sequence. This requires the language model to be integrated to the decoder at an earlier stage of the encoder-decoder model. Gulcehre et al. (2015) introduced a method called deep fusion, in which the hidden states of the language model are concatenated with the hidden states of the decoder (see Figure 2). Like shallow fusion, deep fusion assumes that the ASR and LM models are pretrained, but deep fusion incorporates a learned controller mechanism

$$g_t = \sigma(\mathbf{v}_g^T \mathbf{d}_t^{LM} + b_g) \quad (10)$$

to weight the language model differently depending on the LM hidden state \mathbf{d}_t^{LM} . \mathbf{v}_g^T and b_g are the weight and bias that are learned in the fine tuning and σ is the sigmoid squashing function. During the fine tuning of the integrated system, the context \mathbf{c}_t , decoder hidden state \mathbf{d}_t , and the weighted hidden state of the language model $g_t \mathbf{d}_t^{LM}$ are all concatenated to create the hidden state of the deep fusion layer \mathbf{d}_t^{DF} . The \mathbf{d}_t^{DF} is parametrised with \mathbf{W}_{DF} and \mathbf{b}_{DF} and fed through a softmax layer to generate the conditional distribution output:

$$\mathbf{d}_t^{DF} = [\mathbf{c}_t; \mathbf{d}_t; g_t \mathbf{d}_t^{LM}] \quad (11)$$

$$P(y_t|\mathbf{h}, \mathbf{y}_{<t}) = \operatorname{softmax}(\mathbf{W}_{DF} \mathbf{d}_t^{DF} + \mathbf{b}_{DF}) \quad (12)$$

The weight matrix \mathbf{W}_{DF} and the bias \mathbf{b}_{DF} are also learned in fine tuning. The integrated system is trained while keeping the parameters of the pre-trained ASR and LM models fixed so that the decoder learns to use the hidden states of the LM during the inference of the next word.

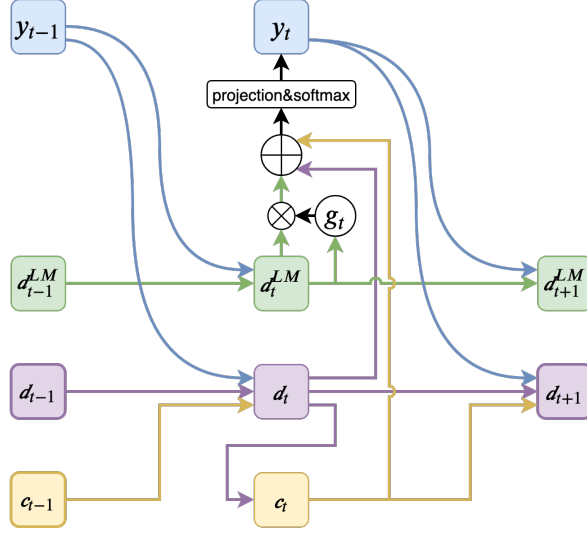


Figure 2: Deep fusion decoding step. Gating function g_t weights the language model state d_t^{LM} , which is concatenated with the context vector c_t and the decoder state d_t .

Deep fusion was originally developed for neural machine translation (NMT), so it is motivated by properties of the translation problem. Gulcehre et al. (2015) give an example when the language model can be more useful in an NMT system: when translating from a language with no articles (a, an, the), such as Chinese, to English, the LM is probably useful in generating the articles because there are no Chinese words that correspond to articles. In contrast, when translating a noun, the LM is probably not as informative. When deep fusion is used in the ASR task, the LM is informative, for instance, when transcribing homophones, such as "two" and "too", and not so informative when transcribing phonetically unambiguous words.

3.3 Cold fusion

In both shallow and deep fusion, the ASR model is initially trained separately from the language model which means that it learns an implicit language model to transcribe the speech. When the external LM is integrated to the system, the implicit LM of the ASR model becomes redundant. This can be viewed as waste of the decoder capacity. Furthermore, the implicitly learned LM can be overfitted, or at least somewhat biased, to the relatively small corpus of transcribed speech. Sriram et al. (2017) developed a method called Cold Fusion, in which the Seq2Seq encoder-decoder is trained to use an external language model from the beginning of the training instead of only at a fine-tuning stage as in deep fusion. This way the Seq2Seq learns to model only the acoustics and rely on the

external LM for language information.

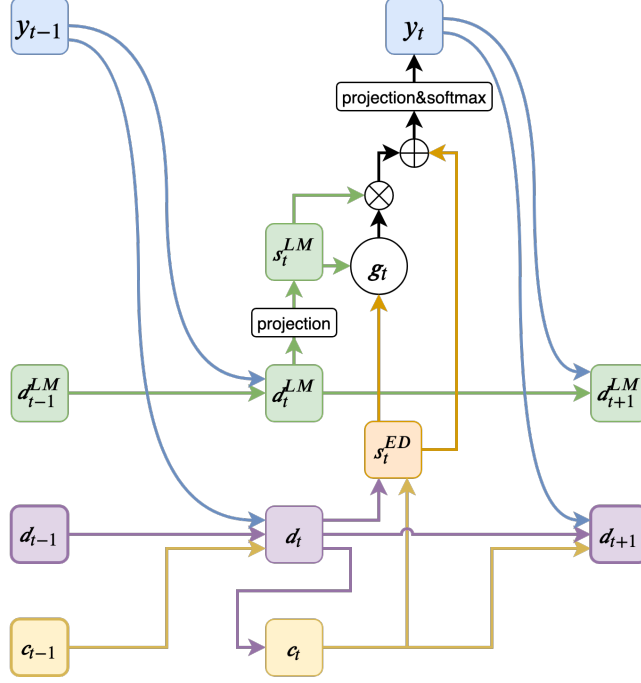


Figure 3: Cold fusion decoding step. The gating function \mathbf{g}_t is a vector, and it depends on both the enc-dec state \mathbf{s}_t^{ED} and the projected language model logit output \mathbf{s}_t^{LM} .

In addition to the early training integration method, there are a number of other differences between deep and cold fusion. In cold fusion, the language model hidden state \mathbf{d}_t^{LM} as used in deep fusion is replaced by the language model logit probability fed through a DNN to create a projection to an embedding space that is uniform across different LMs. The LM hidden state distribution and dynamics can vary considerable across different language models and data, making it impossible to change the LM. This projection enables switching to use a different LM:

$$\mathbf{s}_t^{LM} = DNN(\text{logit } P_{LM}(\mathbf{y})) \quad (13)$$

Furthermore, cold fusion uses the fine-grained gating (Yang et al., 2016) mechanism (equation 15), making \mathbf{g}_t a vector instead of a scalar as in deep fusion:

$$\mathbf{g}_t = \sigma(\mathbf{W}_g[\mathbf{s}_t^{ED}; \mathbf{s}_t^{LM}] + \mathbf{b}_g) \quad (14)$$

$$\mathbf{s}_t^{CF} = [\mathbf{s}_t^{ED}; \mathbf{g}_t \circ \mathbf{s}_t^{LM}] \quad (15)$$

where \circ represents the element-wise (also called Hadamard) product of the vectors. The fact that the gating function \mathbf{g}_t depends both on the LM hidden state \mathbf{s}_t^{LM} and the decoder hidden state \mathbf{s}_t^{ED} makes it possible to weight the LM more in situations where the acoustic signal is noisy, for instance, even though the LM would not have a large weight based only on the LM state as in deep fusion. Furthermore, the fine-grained gating allows for more flexibility in choosing which aspects of the LM state \mathbf{s}_t^{LM} are emphasised at each time step, since the gating \mathbf{g}_t weights each node of the LM state \mathbf{s}_t^{LM} individually.

The cold fusion state \mathbf{s}_t^{CF} goes through the final layers of the decoder to produce the output probability distribution:

$$\mathbf{r}_t^{CF} = \text{DNN}(\mathbf{s}_t^{CF}) \quad (16)$$

$$P(y_t|\mathbf{h}, \mathbf{y}_{<t}) = \text{softmax}(\mathbf{W}_{CF}\mathbf{r}_t^{CF} + \mathbf{b}_{CF}) \quad (17)$$

3.4 Decoder trained on multitask objective

When the decoder input vector is set as a zero context vector, the Seq2Seq decoder objective function reduces to the LM objective function, i.e. the probability distribution of the next output symbol depends only on the previous state of the decoder and not on the encoder output. This allows for training the decoder on a multi-task objective, in which it learns the language modelling task and the usual Seq2Seq decoding task. The objective at each iteration is sampled from the set of two objectives according to some prior weights. When the decoder is trained on the language modelling task, the encoder and attention parameters are not affected.

3.5 LM as a decoder layer

Toshniwal et al. (2018) proposed also using an external language model as a lower decoder layer. The LM layer is placed to the decoder of a pretrained LAS models and the whole model is then fine-tuned on the LAS objective for a few epochs. Similar methods have been proposed by Ramachandran et al. (2016) and Rao et al. (2017) who used a pretrained LM as a lower decoder layer in a machine translation model and as the initialisation weights for an RNN transducer for speech recognition, respectively.

4 Experiments

This section reviews the experiments presented in by Sriram et al. (2017) and Toshniwal et al. (2018). The experiments were designed to determine which LM integration method achieves the best results in terms of word error rate (WER) or character error rate (CER). However, the experiment setups were not uniform, so no clear comparisons can be made.

4.1 Sriram et al. (2017)

4.1.1 Experiment setup

Sriram et al. (2017) performed single-domain and domain transfer experiments. For single domain, they used the LibriSpeech data set of audiobooks (960 hours) and non-overlapping text corpus of 800M tokens (900k unique words) from 14500 books (Panayotov et al., 2015). They also added noise to this data to test the noise robustness of the models. For the domain transfer experiments, they collected a corpus of search queries and a corpus of movie transcripts respectively for source and target domains. Their source data set contained 411k utterances (650 hours), and target dataset 345k utterances (676 hours).

They had three main models that they compared: a baseline attention based encoder-decoder with no external language model, a deep fusion model, and a cold fusion model.

The LM for the single domain experiments trained on the Librispeech data consisted of one layer of 1536-dimensional GRUs. The LM for the domain transfer experiments trained on the search query and movie transcript data consisted of three layers of 1024-dimensional GRUs. They used 40 mel scale filterbank features as input and characters as output.

4.1.2 Results

In the single domain experiments, using clean speech data the baseline attention ASR system got better results than deep fusion, whereas when using noisy data deep fusion got better results. This is in line with the idea that deep fusion enables the decoder to rely more on the language model when the speech is noisy, making this approach more robust to noise. However, cold fusion performed better than shallow or deep fusion on both clean and noisy speech. This can be because the decoder capacity is used more effectively, utilising the external language model better because of the more fine-grained integration.

Because the deep fusion approach uses the language model hidden state \mathbf{d}_t instead of the projected output probabilities as in cold fusion, switching the language model is not possible. This makes domain transfer harder since the language model cannot be simply changed from one trained on the source domain to another trained on the target domain. For this reason they used a single LM for the fusion models. The cold fusion approach got better results than shallow or deep fusion in both source and target domain. The cold fusion approach is more complicated than the other two, including a projection of the LM model state and also a DNN before the final softmax layer after the concatenation of the language model output projection and Seq2Seq model state. This means that cold fusion requires more parameters than the other two approaches, which naturally affects the performance.

Sriram et al. (2017) performed also some ablation studies on the deep and cold fusion approaches. They noted that the projection of the language model output helped especially in the domain transfer tests.

The main reasoning behind training the Seq2Seq model to utilise the LM from the

beginning was that the decoder capacity would be used more efficiently to learn the Seq2Seq task rather than learning an implicit language model. They experimented reducing the size of the decoder, and noted that the cold fusion model was able to perform well with a smaller decoder size compared to deep fusion.

4.2 Toshniwal et al. (2018)

4.2.1 Switchboard data

Toshniwal et al. (2018) used two different speech data types in their experiments. First, they used the Switchboard (SWB) corpus of 300h / 192k utterances of conversational telephone speech and did the evaluation on SWB and CallHome (CH) corpora. They trained the language models on text from the SWB and Fischer corpora with 2M utterances. With these data, they used a model with 4-layer BLSTM 256-dimensional encoder, one-layer LSTM decoder, and one-layer 512-dimensional LSTM language model. They used 40 mel scale filterbank features with deltas and a 1000 BPE wordpiece vocabulary.

4.2.2 Google voice search and dictation data

Second, they used the Google voice search and dictation dataset of 22M utterances. They trained a larger model on this dataset: 5-layer unidirectional LSTM network (1400 dimensions) as encoder, 2-layer unidirectional LSTM (1024 dimensions) as decoder, and 2-layers of 2048-dimensional LSTM cells as the language model. They extracted 80 mel scale filterbank features as the input and had a 16384 BPE wordpiece vocabulary as the output.

4.2.3 Results

In the experiments of Toshniwal et al. (2018), shallow fusion performed the best on the SWB data and on the voice search data. On the dictation data, shallow fusion and cold fusion got similar results. The reason why cold fusion did not achieve the best results on these experiments as it did on the experiments of Sriram et al. (2017), could be that the latter used more parameters in the cold fusion layer than in other fusion approaches.

Toshniwal et al. (2018) evaluated also the methods of using a pretrained LM as a decoder layer or training the decoder on the multi-task objective (LM and Seq2Seq). They got some improvement on the LAS baseline by using each of these methods. The LM as decoder layer approach got better results than the multitask approach. Interestingly, LM as a decoder layer got similar results as the deep fusion and cold fusion approaches in these experiments.

When using a larger LM for rescoring the top-8 hypotheses of the beam search algorithm, the cold fusion approach achieved the best results. It is hard to know exactly why this fusion approach produces better top-8 lists than the other approaches.

5 Conclusion

A few different approaches to integrating an external LM to an E2E ASR system has been recently proposed. The simplest method, shallow fusion, is to combine the output probabilities of the two models by a weighted linear combination in log domain. More complicated approaches, deep and cold fusion, fuse the hidden states of the two models together (or the hidden state of the Seq2Seq model and the projected output of the LM as in cold fusion) to utilise the LM and the Seq2Seq model more effectively, weighting the LM more when it is more useful, weighting different LM nodes differently (fine-grained gating), or utilising the pretrained LM in the Seq2Seq model training.

The more complicated fusing approaches (deep and cold) require additional layers and more stages in the training of the combined models, which makes them more expensive than the shallow fusion approach. However, the more complicated methods do not necessarily get better results than shallow fusion (Toshniwal et al., 2018). This would mean that the shallow fusion is the most cost-effective approach, being the simplest.

Although it is not clear which one of LM integration methods is the best, in general integrating an external language model by any of these methods brings improved results over the baseline LAS model without an external LM, at least with noisy data or in domain transfer tasks (Toshniwal et al., 2018; Sriram et al., 2017; McDermott et al., 2019).

The different fusion methods might be useful in different scenarios. The cold fusion approach is in theory able to utilise the external language model the most since it is fused in the system in the earliest stage, and it is fused with a more fine-grained gating mechanism. This might make the cold fusion most attractive method for domain transfer when there is a parallel corpus in a source domain and a sufficiently large text corpus in the target domain.

References

- Bahdanau, D., Cho, K., and Bengio, Y. (2014). Neural machine translation by jointly learning to align and translate. *arXiv preprint arXiv:1409.0473*.
- Bourlard, H. A. and Morgan, N. (2012). *Connectionist speech recognition: a hybrid approach*, volume 247. Springer Science & Business Media.
- Chan, W., Jaitly, N., Le, Q., and Vinyals, O. (2016). Listen, attend and spell: A neural network for large vocabulary conversational speech recognition. In *2016 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 4960–4964. IEEE.
- Gulcehre, C., Firat, O., Xu, K., Cho, K., Barrault, L., Lin, H.-C., Bougares, F., Schwenk, H., and Bengio, Y. (2015). On using monolingual corpora in neural machine translation. *arXiv preprint arXiv:1503.03535*.

- Jelinek, F. (1976). Continuous speech recognition by statistical methods. *Proceedings of the IEEE*, 64(4):532–556.
- McDermott, E., Sak, H., and Variani, E. (2019). A density ratio approach to language model fusion in end-to-end automatic speech recognition. In *2019 IEEE Automatic Speech Recognition and Understanding Workshop (ASRU)*, pages 434–441. IEEE.
- Panayotov, V., Chen, G., Povey, D., and Khudanpur, S. (2015). Librispeech: an asr corpus based on public domain audio books. In *2015 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 5206–5210. IEEE.
- Ramachandran, P., Liu, P. J., and Le, Q. V. (2016). Unsupervised pretraining for sequence to sequence learning. *arXiv preprint arXiv:1611.02683*.
- Rao, K., Sak, H., and Prabhavalkar, R. (2017). Exploring architectures, data and units for streaming end-to-end speech recognition with rnn-transducer. In *2017 IEEE Automatic Speech Recognition and Understanding Workshop (ASRU)*, pages 193–199. IEEE.
- Sriram, A., Jun, H., Satheesh, S., and Coates, A. (2017). Cold fusion: Training seq2seq models together with language models. *arXiv preprint arXiv:1708.06426*.
- Sutskever, I., Vinyals, O., and Le, Q. V. (2014). Sequence to sequence learning with neural networks. In *Advances in neural information processing systems*, pages 3104–3112.
- Toshniwal, S., Kannan, A., Chiu, C.-C., Wu, Y., Sainath, T. N., and Livescu, K. (2018). A comparison of techniques for language model integration in encoder-decoder speech recognition. In *2018 IEEE Spoken Language Technology Workshop (SLT)*, pages 369–375. IEEE.
- Yang, Z., Dhingra, B., Yuan, Y., Hu, J., Cohen, W. W., and Salakhutdinov, R. (2016). Words or characters? fine-grained gating for reading comprehension. *arXiv preprint arXiv:1611.01724*.